

I590 - Time Series Analysis - Code Portfolio

James Provost

April 21, 2019

Importing / Processing Time Series

Applied Statistical Time Series Analysis - has lots of data sets (astsa).

Reference: <https://www.stat.pitt.edu/stoffer/tsa4/xChanges.htm> for information about astsa.

```
library('astsa')
```

Read in data from a file.

Reference: Manipulating Time Series Data in R with xts & zoo - Chapter 1, Data Camp

```
library(xts)
# Convert either a time series object or zoo object to an xts object
# XTS - eXtensible Time Series - based on a zoo object
my.xts <- as.xts(chicken)

# Write zoo object to file and then read in a previously saved zoo object
write.zoo(my.xts, 'zoo-file.txt')
my.xts2 <- as.xts(read.zoo('zoo-file.txt', FUN = as.yearmon))

# Use regular file reading functions
# read.csv(), read.table(), read.delim() - see R help
```

Create a time series

Specify data, start (can be just a number, or a vector, with the second value referring to frequency), and frequency (number of units in a time period, like quarters or months)

Reference: Introduction to Time Series Analysis - Chapter 1, Data Camp.

```
# Example creates a 60 period time series, from Jan 1980 with 12 periods
# per year. If frequency was 4, then it would start Q1 1980.
my.ts <- ts(seq(1:60), start = c(1980,1), frequency = 12)

my.ts
```

```
##      Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
## 1980   1   2   3   4   5   6   7   8   9  10  11  12
## 1981  13  14  15  16  17  18  19  20  21  22  23  24
## 1982  25  26  27  28  29  30  31  32  33  34  35  36
## 1983  37  38  39  40  41  42  43  44  45  46  47  48
## 1984  49  50  51  52  53  54  55  56  57  58  59  60
```

```
my.ts <- ts(seq(1:60), start = c(1980,1), frequency = 4)
```

```
my.ts
```

```
##      Qtr1 Qtr2 Qtr3 Qtr4
## 1980    1    2    3    4
## 1981    5    6    7    8
## 1982    9   10   11   12
```

```
## 1983  13  14  15  16
## 1984  17  18  19  20
## 1985  21  22  23  24
## 1986  25  26  27  28
## 1987  29  30  31  32
## 1988  33  34  35  36
## 1989  37  38  39  40
## 1990  41  42  43  44
## 1991  45  46  47  48
## 1992  49  50  51  52
## 1993  53  54  55  56
## 1994  57  58  59  60
```

Create an xts object.

Reference: Introduction to Time Series Analysis - Chapter 1, Data Camp.

```
library(xts)
# Build a sample matrix and index vector of dates
my.matrix <- matrix(1:5, ncol = 1, nrow = 5)
my.index <- as.Date(c('2010-01-01', '2011-01-01', '2012-01-01', '2013-01-01', '2014-01-01'))
# Create an xts object
my.xts <- xts(my.matrix, order.by = my.index)
my.xts
```

```
##           [,1]
## 2010-01-01    1
## 2011-01-01    2
## 2012-01-01    3
## 2013-01-01    4
## 2014-01-01    5
```

Basic exploration

Basic functions to evaluate aspects of a time series object.

Reference: Introduction to Time Series Analysis, Data Camp.

```
# Period the time series starts
start(my.ts)
```

```
## [1] 1980    1
```

```
# Period it ends
end(my.ts)
```

```
## [1] 1994    4
```

```
# Frequency
frequency(my.ts)
```

```
## [1] 4
```

```
# The interval from one period to another in terms of time units (1/frequency)
deltat(my.ts)
```

```
## [1] 0.25
```

```
# Whether it's a time series object
is.ts(my.ts)
```

```
## [1] TRUE
```

```
# Vector of indices
```

```
time(my.ts)
```

```
##      Qtr1    Qtr2    Qtr3    Qtr4
## 1980 1980.00 1980.25 1980.50 1980.75
## 1981 1981.00 1981.25 1981.50 1981.75
## 1982 1982.00 1982.25 1982.50 1982.75
## 1983 1983.00 1983.25 1983.50 1983.75
## 1984 1984.00 1984.25 1984.50 1984.75
## 1985 1985.00 1985.25 1985.50 1985.75
## 1986 1986.00 1986.25 1986.50 1986.75
## 1987 1987.00 1987.25 1987.50 1987.75
## 1988 1988.00 1988.25 1988.50 1988.75
## 1989 1989.00 1989.25 1989.50 1989.75
## 1990 1990.00 1990.25 1990.50 1990.75
## 1991 1991.00 1991.25 1991.50 1991.75
## 1992 1992.00 1992.25 1992.50 1992.75
## 1993 1993.00 1993.25 1993.50 1993.75
## 1994 1994.00 1994.25 1994.50 1994.75
```

```
# Position in cycle of the observation
```

```
cycle(my.ts)
```

```
##      Qtr1 Qtr2 Qtr3 Qtr4
## 1980     1     2     3     4
## 1981     1     2     3     4
## 1982     1     2     3     4
## 1983     1     2     3     4
## 1984     1     2     3     4
## 1985     1     2     3     4
## 1986     1     2     3     4
## 1987     1     2     3     4
## 1988     1     2     3     4
## 1989     1     2     3     4
## 1990     1     2     3     4
## 1991     1     2     3     4
## 1992     1     2     3     4
## 1993     1     2     3     4
## 1994     1     2     3     4
```

```
# Pulls out part of a TS between specified start and end periods
```

```
window(my.ts, start = c(1980,7), end = c(1980,12))
```

```
##      Qtr1 Qtr2 Qtr3 Qtr4
## 1981           7     8
## 1982     9    10    11    12
```

Data Manipulation

Aggregation nfrequency new number of observations per unit of time; must be a divisor of the frequency of x.
FUN aggregation function

Reference: Metcalfe, A. and Cowpertwait, P. (2009). *Introductory Time Series with R*. New York, NY; Spring-Veriag, New York, p. 17

```
library(astsa)
# Sum by quarter
aggregate(chicken, nfrequency = 4, FUN = sum)
```

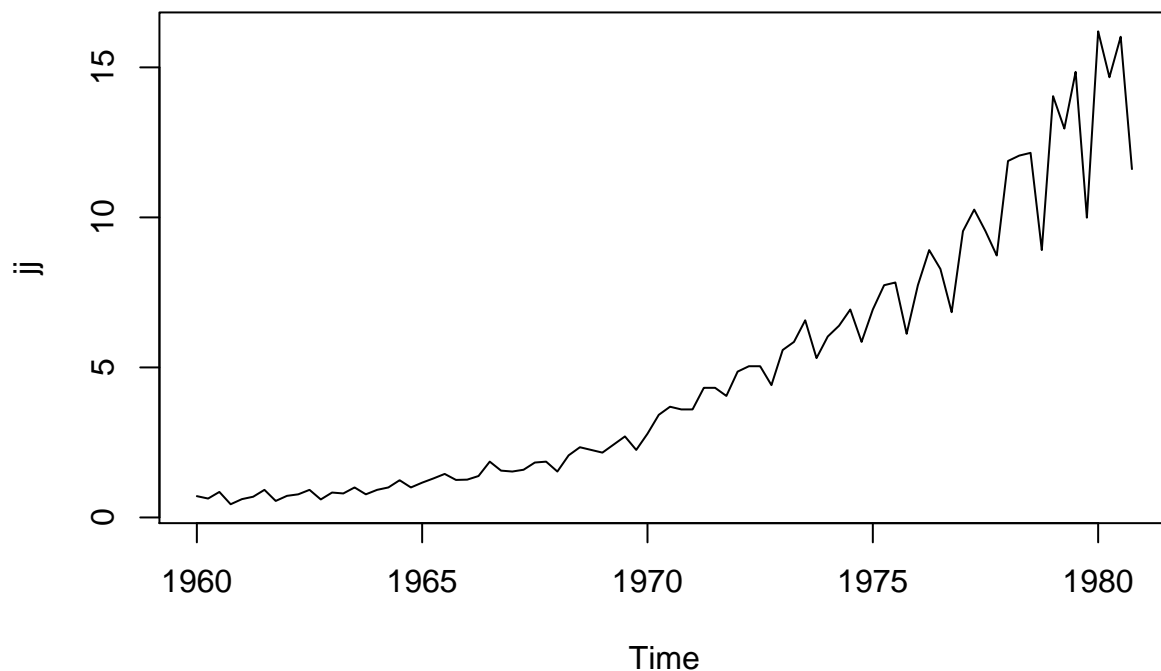
```
## Time Series:
## Start = 2001.583333333333
## End = 2016.333333333333
## Frequency = 4
## [1] 197.76 190.50 188.15 191.36 191.15 185.66 190.52 198.09 206.30 206.92
## [11] 219.43 237.16 234.26 221.75 221.63 223.52 224.34 213.76 204.99 206.75
## [21] 210.76 209.68 228.24 241.44 242.57 231.87 242.42 257.42 264.36 261.20
## [31] 257.81 263.69 254.50 247.32 252.58 261.46 262.27 255.81 257.55 260.82
## [41] 266.11 269.40 277.21 283.29 286.51 293.41 304.19 316.05 317.76 313.34
## [51] 317.02 332.27 340.20 341.80 342.92 347.86 344.73 338.94 335.21 335.28
```

Stationarity - Stable - Mean remains constant - there is no trend - Correlation from period to period remains constant

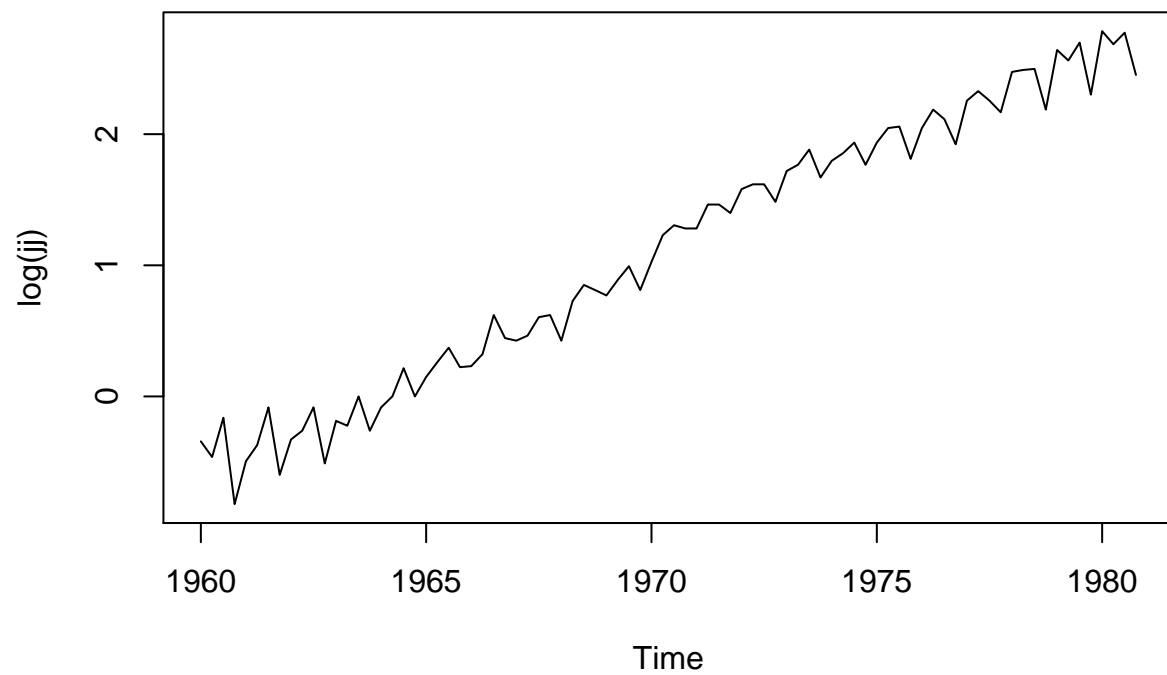
Data Transformation. The `diff()` function shows the difference from period to period in the time series. It is a way to remove the trend (including Random Walk) from a time series. The `log()` function will help remove a growth in variability over time (like a multiplicative trend). If you `diff(log())`, you could make this type of data stationary.

Reference: ARIMA Modeling with R - Chapter 1, Data Camp.

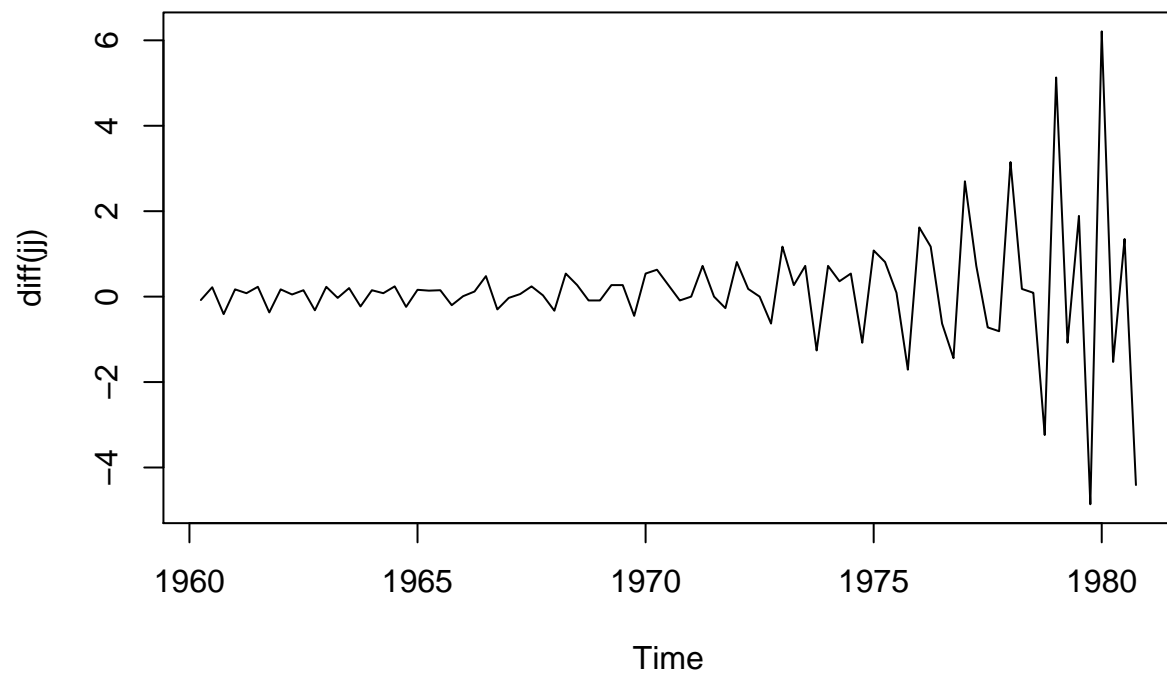
```
library(astsa)
# Using Johnson & Johnson quarterly earnings, because it has a growing variability in the trend
ts.plot(jj)
```



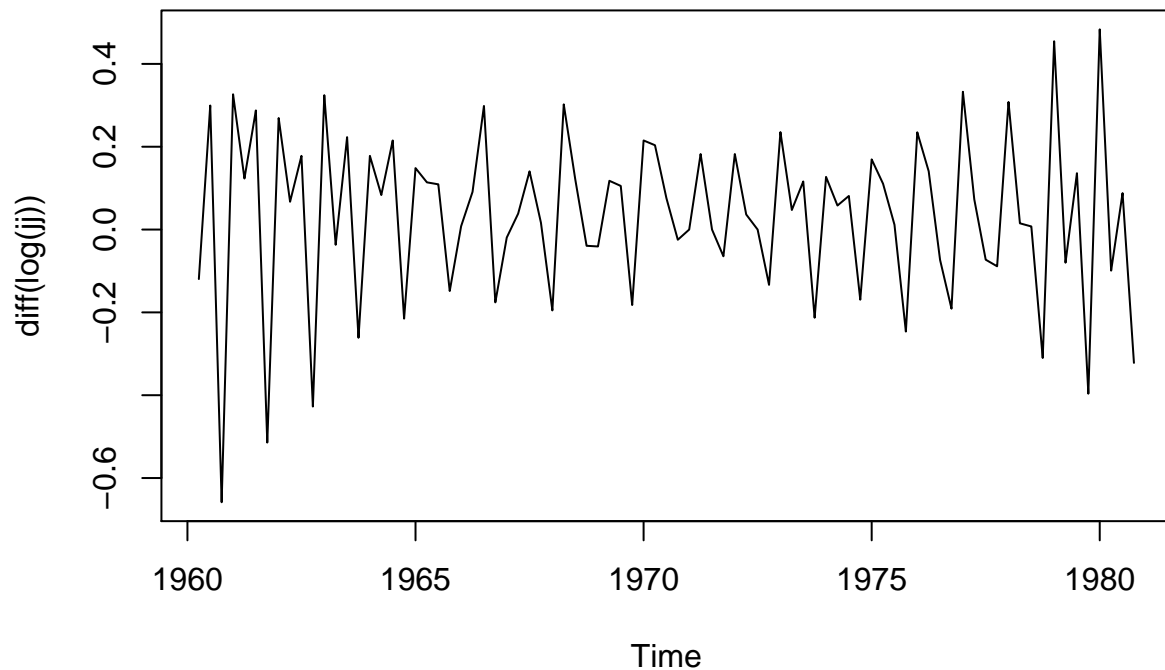
```
# Take out the upward curve, make it a more linear trend  
ts.plot(log(jj))
```



```
# Take out the trend alone  
ts.plot(diff(jj))
```



```
# Take out the curve then take out the trend  
ts.plot(diff(log(jj)))
```



Box-Cox Transformations. A transformation for stabilizing variance, usually from -1 (inverse transformation) to 1 (no transformation), with things like natural log and square root in between. `BoxCox.lambda()` will determine the best lambda value.

Reference: Forecasting Using R - Chapter 4, Data Camp.

```
library(forecast)
BoxCox.lambda(chicken)
```

```
## [1] 1.686301
```

Exploratory Visualization of Time Series

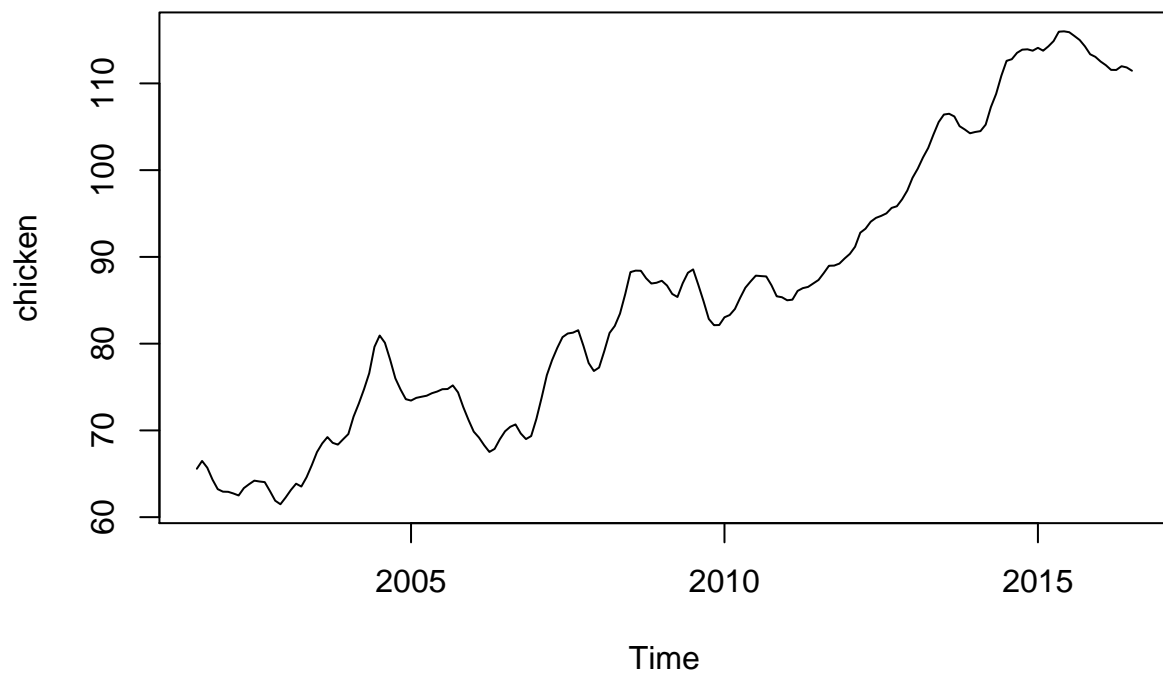
Basic plotting

References: Forecasting Using R, Chapter 1, Data Camp. <https://www.rstudio.com/wp-content/uploads/2015/03/ggplot2-cheatsheet.pdf> https://cran.r-project.org/web/packages/ggfortify/vignettes/plot_ts.html

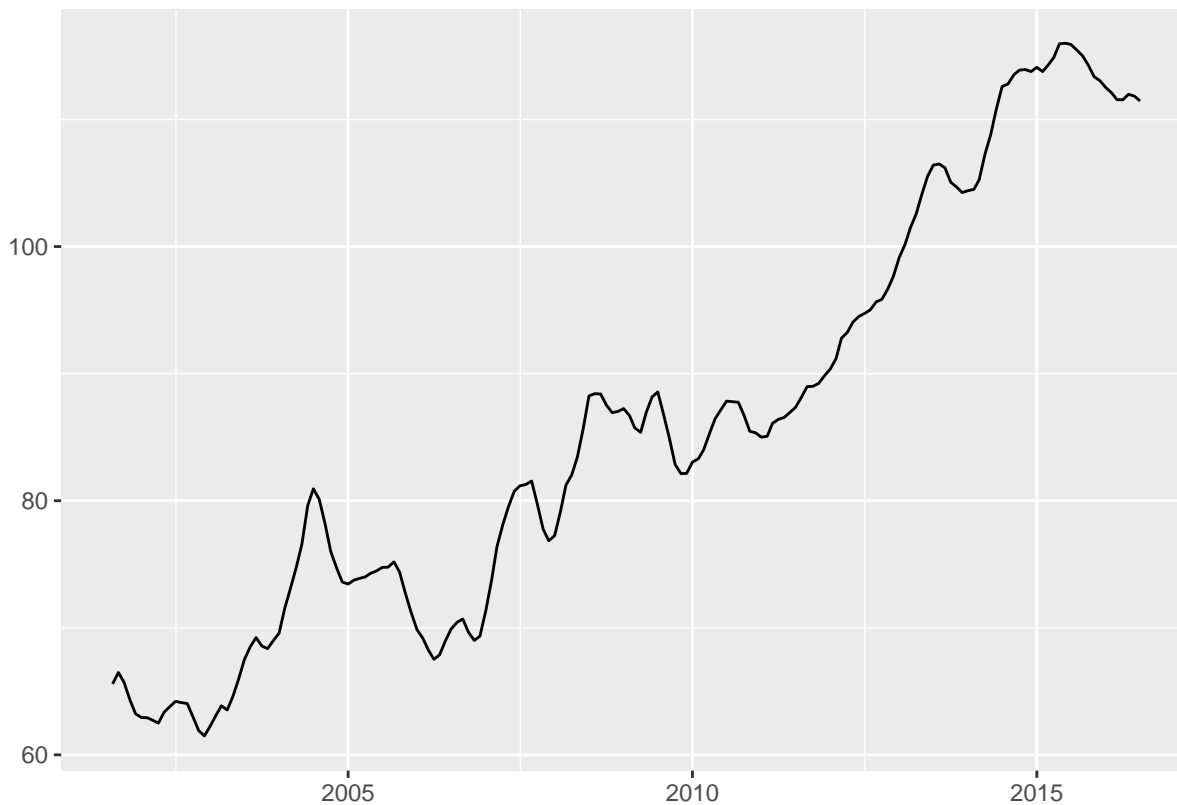
```
library(astsa)
library(ggplot2)

plot(chicken)

ts.plot(chicken)
```



```
# Library ggfortify is needed for fortify to handle time series objects  
library(ggfortify)  
autoplot(chicken)
```

Also refer to `decompose()`, `sarima()`, `checkresiduals()`, `acf()`, `pacf()` and `acf2()` functions in the Time Series Analysis section.

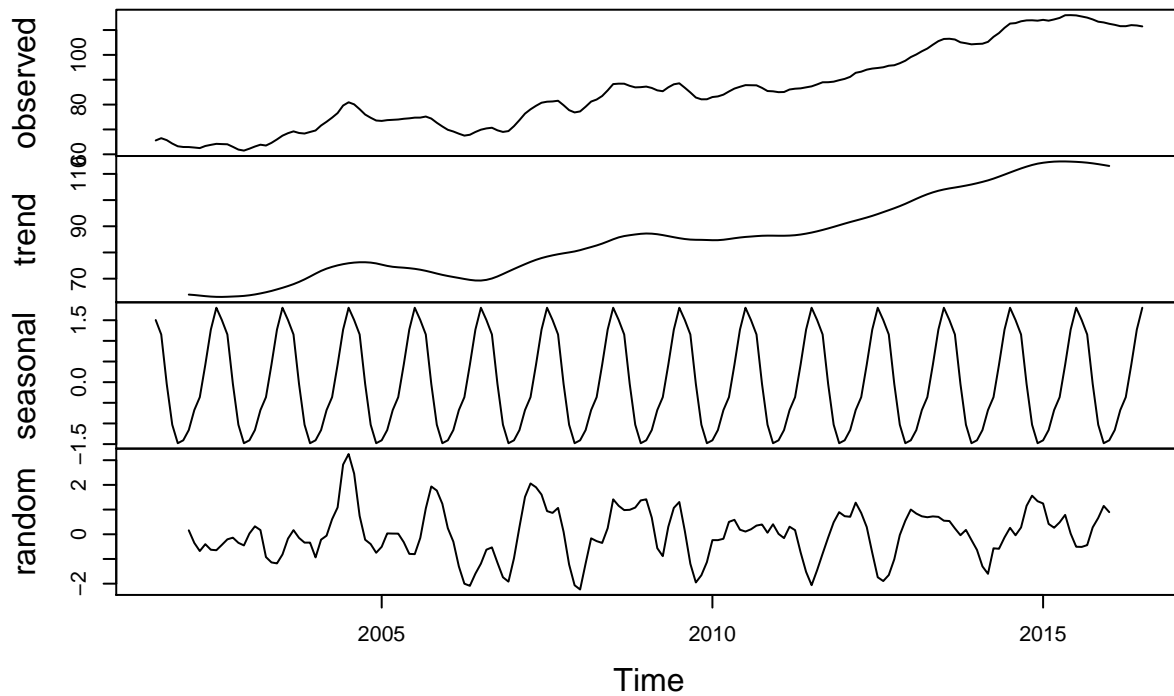
Time Series Analysis

Decompose - Split a TS into components for Trend, Seasonal and Random (Residual)

Reference: Metcalfe, A. and Cowpertwait, P. (2009). *Introductory Time Series with R*. New York, NY; Springer-Verlag, New York, p. 22

```
plot(decompose(chicken))
```

Decomposition of additive time series

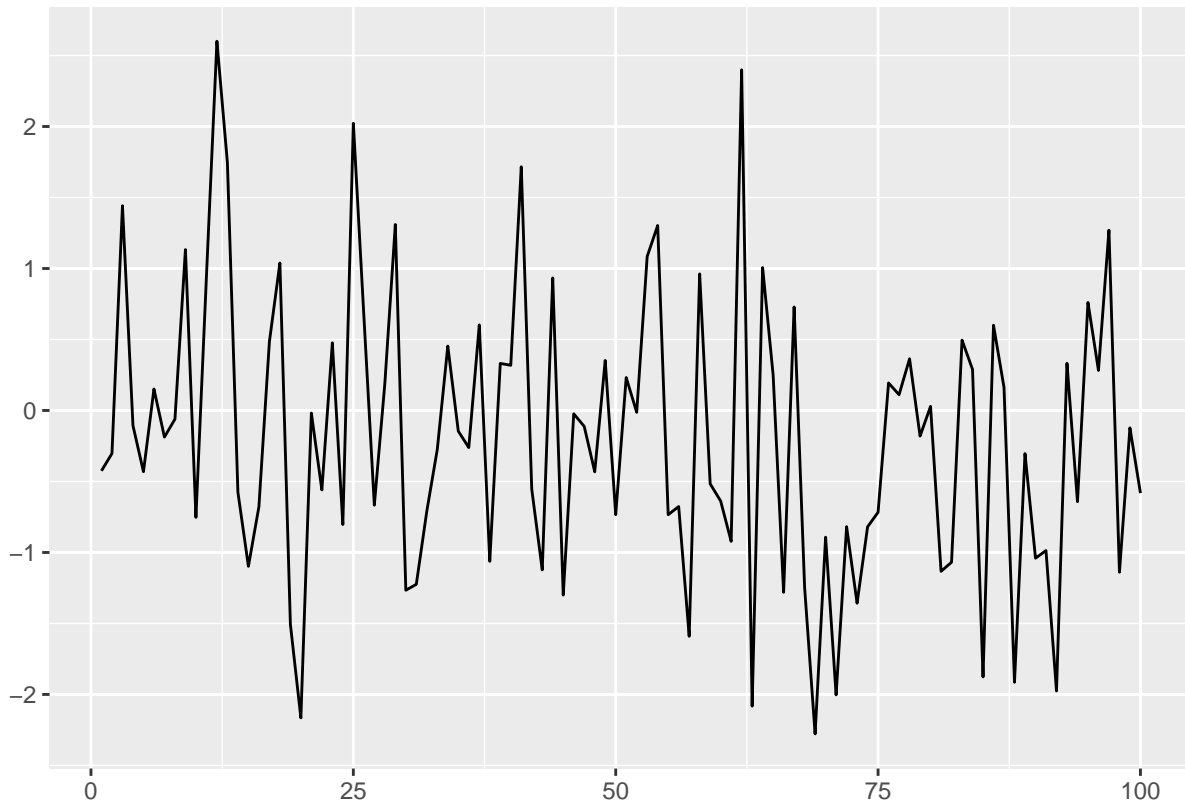


ARIMA - Auto Regressive Integrated Moving Average

To build an ARIMA time series use `arima.sim()`. Specify the level of each of AR, Differencing and MA. Optionally, specify the coefficients in additional parameters in the list, like `ma` or `ar`.

Reference: ARIMA Modeling with R - Chapter 1, Data Camp.

```
# Example of a White Noise model: 0 AR, 0 Diff, 0 MA
my.wn <- arima.sim(model = list(order = c(0,0,0)), n = 100)
autoplot(my.wn)
```



To estimate an ARIMA model, use the `arima()` function to build such as model and then evaluate.

This is some code to cycle through a bunch of models to see which one has the best AIC (you could use a different criterion, like BIC)

Reference: Metcalfe, A. and Cowpertwait, P. (2009). *Introductory Time Series with R*. New York, NY; Springer-Verlag, New York, p. 131

```
# Set the default - white noise
best.order <- c(0, 0, 0)
# Initialize the AIC score
best.aic <- Inf
# This is just looping through a few AR and MA models, may want to use
# acf() and pacf() to get a sense for the size of these loops
for (i in 0:2) for (j in 0:2) {
  # Calculate the AIC for the next type of ARIMA model
  fit.aic <- AIC(arima(diff(log(chicken)), order = c(i, 0, j)))
  # If the new AIC is lowest, keep it, the type of model and the model
  if (fit.aic < best.aic) {
    best.order <- c(i, 0, j)
    best.arma <- arima(diff(log(chicken)), order = best.order)
    best.aic <- fit.aic
  }
}

# Display the best type of model
best.order
```

```
## [1] 2 0 1
```

Can also build models using the `sarima()` function, supplying the AR, Differencing and MA parameters. This function will also build several visualizations of the residuals of the model to help evaluate if this is a good model or not.

Reference: ARIMA Modeling with R - Chapter 2, Data Camp.

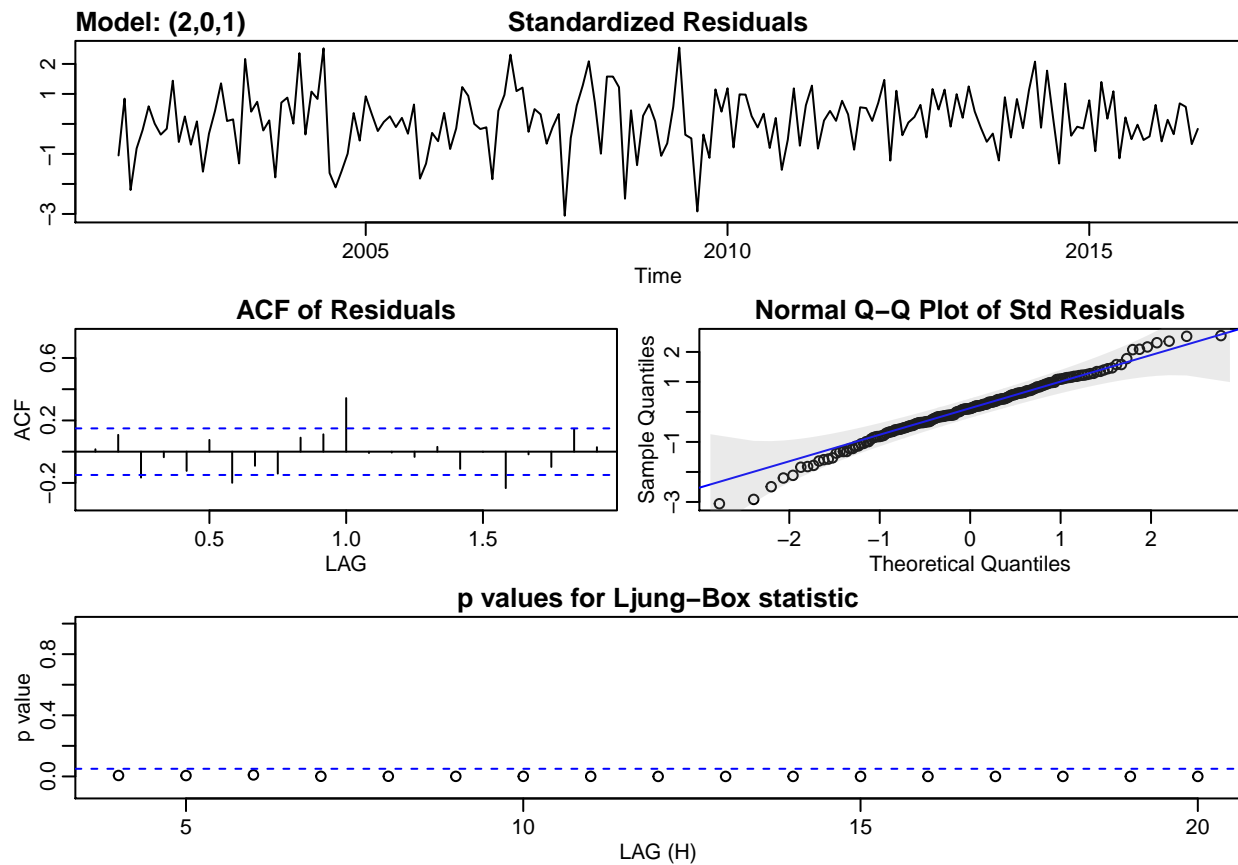
```
library(astsa)
my.arima <- sarima(chicken, p = 2, d = 0, q = 1)
```

```
## initial   value 2.784868
## iter    2 value 2.197149
## iter    3 value 2.180068
## iter    4 value 1.403829
## iter    5 value 0.907290
## iter    6 value -0.070224
## iter    7 value -0.095722
## iter    8 value -0.108832
## iter    9 value -0.213591
## iter   10 value -0.213600
## iter   11 value -0.271682
## iter   12 value -0.333618
## iter   13 value -0.383427
## iter   14 value -0.399481
## iter   15 value -0.400894
## iter   16 value -0.400975
## iter   17 value -0.401010
## iter   18 value -0.401031
## iter   19 value -0.401032
## iter   20 value -0.401045
## iter   21 value -0.401956
## iter   22 value -0.402457
## iter   23 value -0.403025
## iter   24 value -0.403329
## iter   25 value -0.404528
## iter   26 value -0.405032
## iter   27 value -0.405785
## iter   28 value -0.405808
## iter   29 value -0.405816
## iter   30 value -0.405824
## iter   31 value -0.405825
## iter   32 value -0.405827
## iter   33 value -0.405829
## iter   34 value -0.405838
## iter   35 value -0.405855
## iter   36 value -0.405897
## iter   37 value -0.405898
## iter   38 value -0.405930
## iter   39 value -0.405943
## iter   40 value -0.405944
## iter   41 value -0.405944
## iter   42 value -0.405946
## iter   43 value -0.405951
## iter   44 value -0.405963
## iter   45 value -0.405990
```

```
## iter 46 value -0.405991
## iter 47 value -0.406014
## iter 48 value -0.406015
## iter 49 value -0.406016
## iter 50 value -0.406016
## iter 51 value -0.406017
## iter 52 value -0.406020
## iter 53 value -0.406026
## iter 54 value -0.406039
## iter 55 value -0.406040
## iter 56 value -0.406051
## iter 57 value -0.406055
## iter 58 value -0.406056
## iter 59 value -0.406056
## iter 60 value -0.406058
## iter 61 value -0.406063
## iter 62 value -0.406075
## iter 63 value -0.406098
## iter 64 value -0.406098
## iter 65 value -0.406116
## iter 66 value -0.406118
## iter 67 value -0.406118
## iter 68 value -0.406118
## iter 69 value -0.406121
## iter 70 value -0.406124
## iter 71 value -0.406133
## iter 72 value -0.406149
## iter 73 value -0.406149
## iter 74 value -0.406153
## iter 75 value -0.406161
## iter 76 value -0.406162
## iter 77 value -0.406162
## iter 78 value -0.406162
## iter 79 value -0.406162
## iter 80 value -0.406162
## iter 81 value -0.406163
## iter 81 value -0.406163
## final value -0.406163
## converged
## initial value -0.366038
## iter 2 value -0.367487
## iter 3 value -0.369377
## iter 4 value -0.373230
## iter 5 value -0.374092
## iter 6 value -0.374742
## iter 7 value -0.374976
## iter 8 value -0.375007
## iter 9 value -0.375011
## iter 10 value -0.375022
## iter 11 value -0.375027
## iter 12 value -0.375042
## iter 13 value -0.375050
## iter 14 value -0.375121
## iter 15 value -0.375276
```

```
## iter 16 value -0.375750
## iter 17 value -0.376215
## iter 18 value -0.376781
## iter 19 value -0.377474
## iter 20 value -0.377820
## iter 21 value -0.377997
## iter 22 value -0.378079
## iter 23 value -0.378087
## iter 24 value -0.378323
## iter 25 value -0.378397
## iter 26 value -0.378424
## iter 27 value -0.378432
## iter 28 value -0.378456
## iter 29 value -0.378466
## iter 30 value -0.378476
## iter 31 value -0.378600
## iter 32 value -0.378659
## iter 33 value -0.378758
## iter 34 value -0.378798
## iter 35 value -0.378815
## iter 36 value -0.378819
## iter 37 value -0.378843
## iter 38 value -0.378852
## iter 39 value -0.378863
## iter 40 value -0.378869
## iter 41 value -0.378878
## iter 42 value -0.378902
## iter 43 value -0.378943
## iter 44 value -0.378994
## iter 45 value -0.379022
## iter 46 value -0.379029
## iter 47 value -0.379032
## iter 48 value -0.379125
## iter 49 value -0.379127
## iter 50 value -0.379128
## iter 51 value -0.379130
## iter 52 value -0.379132
## iter 53 value -0.379136
## iter 54 value -0.379144
## iter 55 value -0.379159
## iter 56 value -0.379167
## iter 57 value -0.379173
## iter 58 value -0.379174
## iter 59 value -0.379174
## iter 60 value -0.379176
## iter 61 value -0.379178
## iter 62 value -0.379181
## iter 63 value -0.379182
## iter 64 value -0.379182
## iter 65 value -0.379182
## iter 66 value -0.379186
## iter 67 value -0.379186
## iter 68 value -0.379186
## iter 69 value -0.379187
```

```
## iter 70 value -0.379187
## iter 71 value -0.379188
## iter 72 value -0.379189
## iter 73 value -0.379189
## iter 74 value -0.379190
## iter 75 value -0.379190
## iter 76 value -0.379190
## iter 76 value -0.379190
## iter 76 value -0.379190
## final value -0.379190
## converged
```

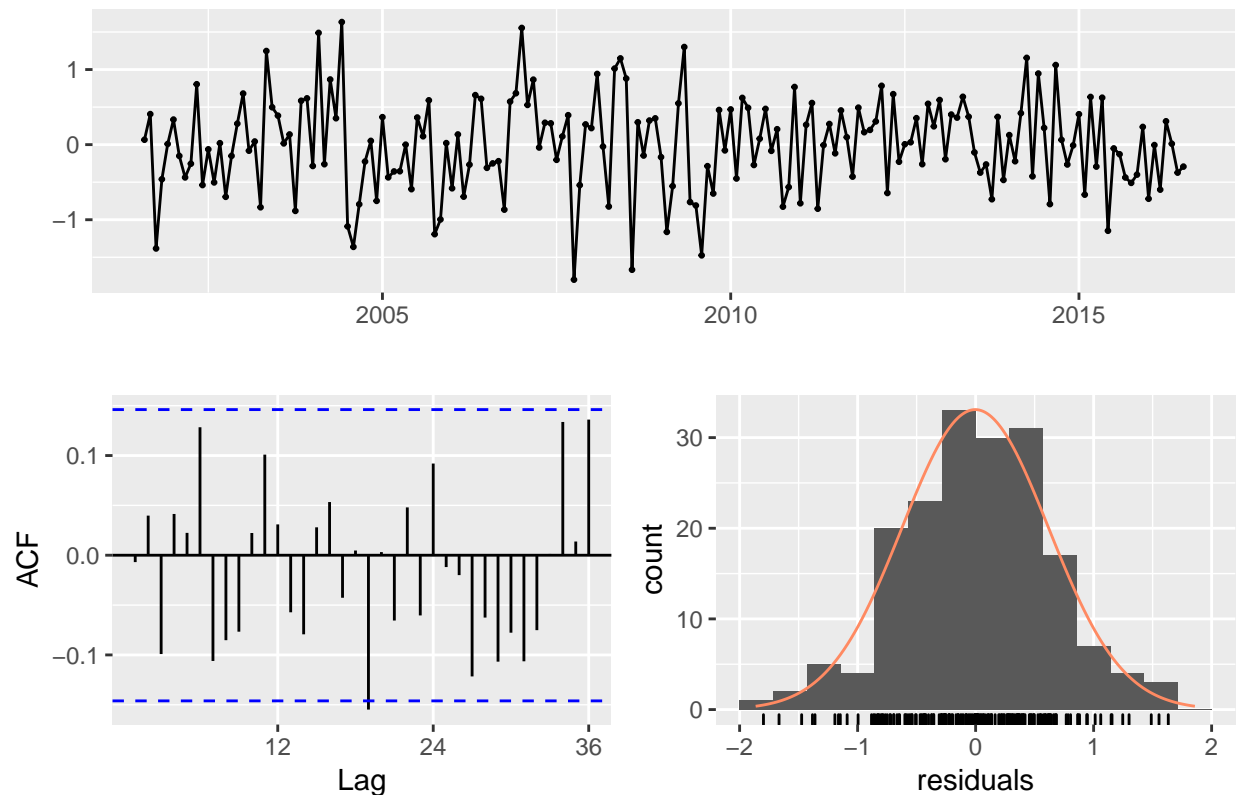


`checkresiduals()` does analysis of a model's residuals, similar to `sarima()` earlier.

Reference: Forecasting Using R - Chapter 2, Data Camp.

```
library(forecast)
checkresiduals(auto.arima(chicken))
```

Residuals from ARIMA(2,1,1)(0,0,1)[12] with drift



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(2,1,1)(0,0,1)[12] with drift
## Q* = 24.287, df = 19, p-value = 0.1854
##
## Model df: 5.   Total lags used: 24
```

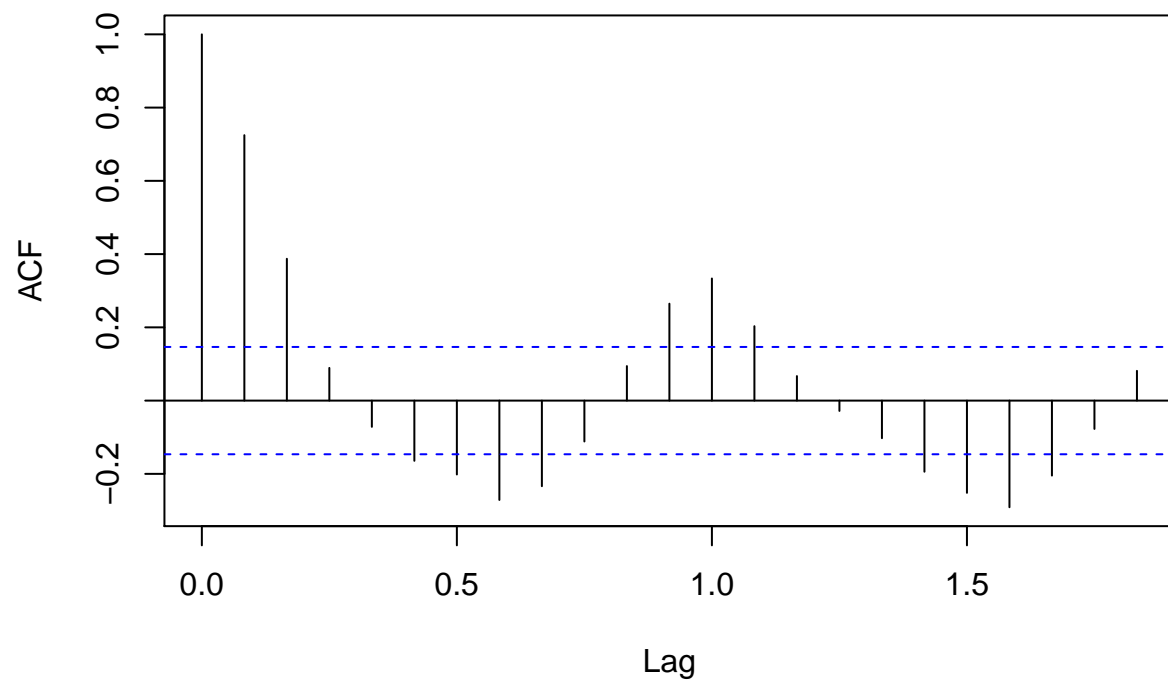
To evaluate the model you might use, use the auto correlation function `acf()` and the partial auto correlation function `pacf()`.

AR model: ACF tails off, PACF cuts off at lag MA model: ACF cuts off at lag, PACF tails off ARMA model: both tail off

Reference: ARIMA Modeling with R - Chapter 2 and Chapter 3, Data Camp.

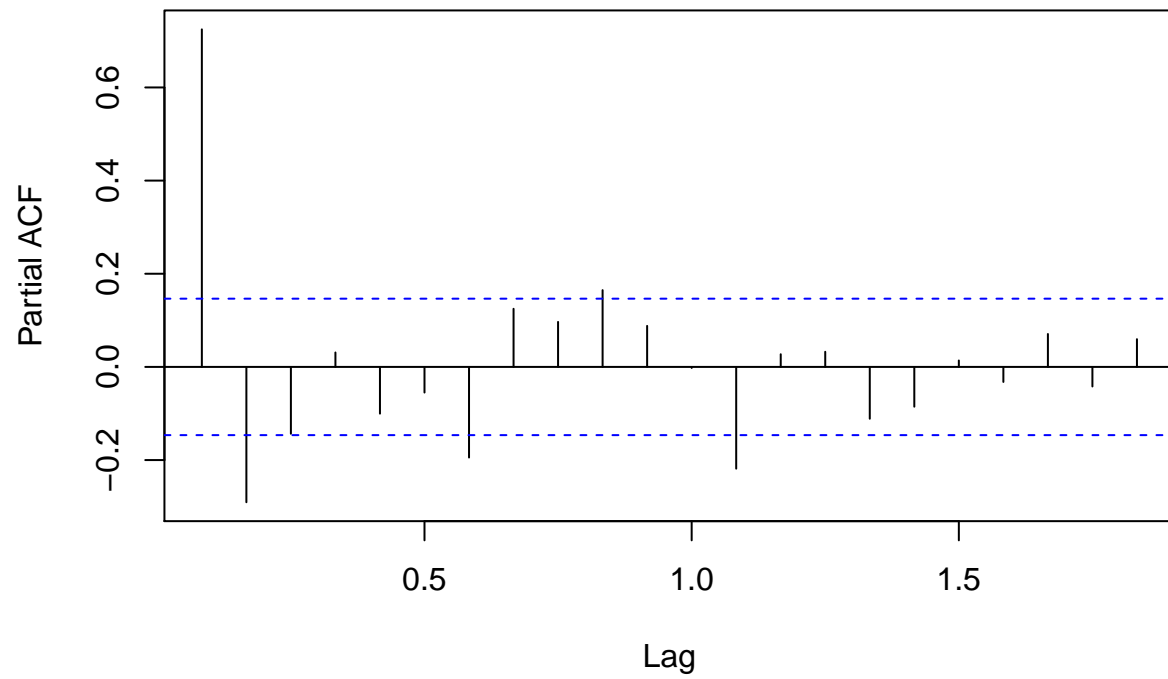
```
acf(diff(chicken))
```


Series diff(chicken)

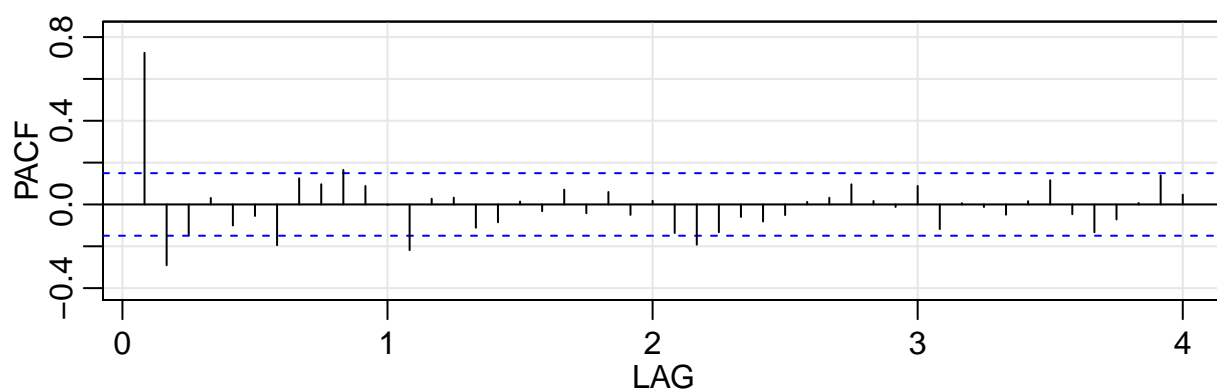
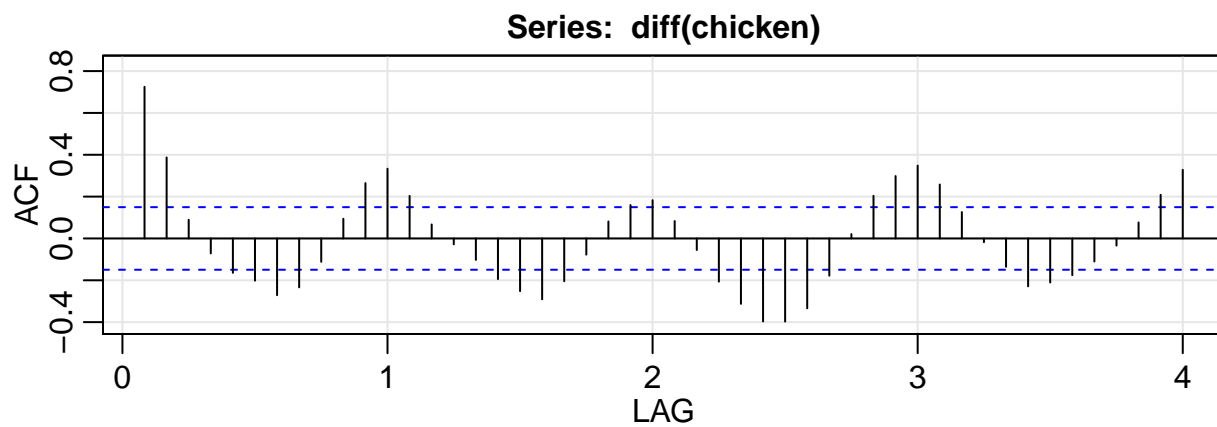


```
pacf(diff(chicken))
```

Series diff(chicken)



```
# Alternatively  
acf2(diff(chicken))
```



##		ACF	PACF
##	[1,]	0.72	0.72
##	[2,]	0.39	-0.29
##	[3,]	0.09	-0.14
##	[4,]	-0.07	0.03
##	[5,]	-0.16	-0.10
##	[6,]	-0.20	-0.06
##	[7,]	-0.27	-0.19
##	[8,]	-0.23	0.12
##	[9,]	-0.11	0.10
##	[10,]	0.09	0.16
##	[11,]	0.26	0.09
##	[12,]	0.33	0.00
##	[13,]	0.20	-0.22
##	[14,]	0.07	0.03
##	[15,]	-0.03	0.03
##	[16,]	-0.10	-0.11
##	[17,]	-0.19	-0.09
##	[18,]	-0.25	0.01
##	[19,]	-0.29	-0.03
##	[20,]	-0.20	0.07
##	[21,]	-0.08	-0.04
##	[22,]	0.08	0.06
##	[23,]	0.16	-0.05
##	[24,]	0.18	0.02
##	[25,]	0.08	-0.14

```
## [26,] -0.06 -0.19
## [27,] -0.21 -0.13
## [28,] -0.31 -0.06
## [29,] -0.40 -0.08
## [30,] -0.40 -0.05
## [31,] -0.33  0.01
## [32,] -0.18  0.03
## [33,]  0.02  0.10
## [34,]  0.20  0.02
## [35,]  0.30 -0.01
## [36,]  0.35  0.09
## [37,]  0.26 -0.12
## [38,]  0.13  0.01
## [39,] -0.02 -0.01
## [40,] -0.14 -0.05
## [41,] -0.23  0.02
## [42,] -0.21  0.12
## [43,] -0.18 -0.05
## [44,] -0.11 -0.13
## [45,] -0.03 -0.07
## [46,]  0.08  0.01
## [47,]  0.21  0.14
## [48,]  0.33  0.05
```

Error Trend and Seasonality models. `ets()` function picks the best model using AICc. Get the lambda from `BoxCox.lambda()` function (mentioned earlier).

Reference: Forecasting Using R - Chapter 4, Data Camp.

```
ets(chicken, lambda = 1.686)
```

```
## ETS(A,Ad,N)
##
## Call:
## ets(y = chicken, lambda = 1.686)
##
## Box-Cox transformation: lambda= 1.686
##
## Smoothing parameters:
##   alpha = 0.9999
##   beta  = 0.9997
##   phi   = 0.8
##
## Initial states:
##   l = 696.4935
##   b = -3.8061
##
## sigma: 14.6426
##
##      AIC      AICc      BIC
## 1907.877 1908.363 1927.035
```

`auto.arima()` can also be used to pick a model and can handle seasonality.

Reference: Forecasting Using R - Chapter 4, Data Camp.

```

auto.arima(chicken)

## Series: chicken
## ARIMA(2,1,1)(0,0,1)[12] with drift
##
## Coefficients:
##          ar1      ar2      ma1      sma1      drift
##          1.2933 -0.5375 -0.4019  0.2756  0.2518
## s.e.  0.2220   0.1542   0.2569  0.0692  0.1428
##
## sigma^2 estimated as 0.396:  log likelihood=-169.51
## AIC=351.01   AICc=351.5   BIC=370.14

```